

Meltdown and Spectre Samples

Written in Assembly

U. Plonus

March 12, 2018

Contents

1	Introduction	5
1.1	Overview	5
1.2	Nasm	5
1.3	Conventions	5
1.3.1	Introduction	5
1.3.2	Data Sections	5
2	Cache Access Timing	7
2.1	Introduction	7
2.2	Detect Cache Access Time	7
2.2.1	High Resolution Timer	7
2.2.2	Cache Access Time Routine	8
2.3	Measure Cache Access Time	8
2.3.1	Setup	8
2.3.2	Measure Time	9
2.4	Read Array via Cache Access Time	12
3	Signals	15
3.1	Basics	15
3.2	Detecting Signals	15
3.3	Handling Signals	15
4	Utilities	17
4.1	Introduction	17
4.2	Exit Program	17
4.3	Random Number Generator	18
4.4	Printing Strings	19
4.4.1	Printing Strings with Length	19
4.4.2	Printing C-Strings	19
4.5	Printing Numbers	21
4.5.1	Printing a 64bit Unsigned Integer	21
A	Glossary	25
B	Acronyms	27
C	x86-Instructions	29

1 Introduction

1.1 Overview

TBD

1.2 Nasm

TBD

```
5  <preamble 5>≡ (12a 14)
    bits 64
        global      _start
        pagesize    equ 4096
```

Defines:

`_start`, used in chunks 9c and 13a.

`pagesize`, used in chunks 8b, 9e, 12, and 13.

1.3 Conventions

1.3.1 Introduction

In this section we define some convention that are specific for this document.

1.3.2 Data Sections

The data is divided into three parts: read-only data, initialized data and uninitialized data. Code chunks with this type of data will all have defined sufficies.

Definition 1 *Read-only data is data that is not modified during program execution. The suffix for read-only data is **-rodata**.*

Definition 2 *Initialized data is data that is changeable during program execution. The data is already initialized with data when the program starts. The suffix for initialized data is **-idata**.*

Definition 3 *Uninitialized data is data that is changeable during program execution. The data is not initialized. The suffix for uninitialized data is **-udata**.*

2 Cache Access Timing

2.1 Introduction

TBD

2.2 Detect Cache Access Time

2.2.1 High Resolution Timer

First we need a high resolution timer to determine the cache access time. For this we use the time stamp counter. The time stamp counter is monotonically incrementing. When reading the time stamp counter (with `rdtsc`) the result is delivered back in the registers `EDX` and `EAX` forming a 64bit value. The time stamp counter is not an absolute value but a relative value, meaning that you cannot (easily) calculate from the time stamp counter to some time units (e.g. ns). But this is no problem as we only want to measure relative times.

To retrieve a 64bit value for the time we shift the value in `EDX` 32 bits to the left and add the value of `EAX` to this.

7 $\langle tsc-64bit \ 7 \rangle \equiv$ (8a)

```
rdtsc
shl    RDX,32
add    RAX,RDX
```

2.2.2 Cache Access Time Routine

Next we need a routine that calculates the cache access time for us.

First we have to ensure in this routine that the speculative execution of the processor does not interfere with our time measurement. For this we use the instruction `lfence` which ensures that all previous reads are done before executing the next instructions.

Next we access a memory location with the address `RDI` by loading this into `RCX` and measure the time before and after the access.

The command `lfence` before reading the time stamp counter is needed because we have to ensure that all reads before the time measurements are done.

At last we calculate the relative time needed to access the memory location. In theory we should see a difference whether the memory location is accessed before or not.

Parameters

`RDI` the address of the memory which is loaded either from the cache or from memory

```
8a  <calculate-cache-access-time 8a>≡ (12a 14)
    _calccachetime:
        lfence
        <tsc-64bit 7>
        mov     R8,RAX
        mov     RCX,[RDI]
        lfence
        <tsc-64bit 7>
        sub     RAX,R8
        ret
```

Defines:

`_calccachetime`, used in chunks 10b and 11a.

2.3 Measure Cache Access Time

2.3.1 Setup

To measure the cache timing we create a standalone program that shows us the time for a cached and for an uncached memory access.

First we need some area in memory with data which we can later read from. This data area goes into the area `.bss` which contains uninitialized data. We align the data at a page boundary and reserve 2 pages for our data.

```
8b  <cachetiming-udata 8b>≡ (12a) 9b▷
        align     pagesize
        data:     resb pagesize
```

Defines:

`data`, used in chunks 9 and 11–13.

Uses `pagesize` 5.

From time to time we need a small scratch area so we define an area with 32 bytes.

9a $\langle \text{scratch-data } 9a \rangle \equiv$ (9b)
 scratch: **resb 32**

Defines:

scratch, used in chunks 10e and 11c.

9b $\langle \text{cachetiming-udata } 8b \rangle + \equiv$ (12a) $\langle 8b$
 $\langle \text{scratch-data } 9a \rangle$

The program begins with the label **_start**.

9c $\langle \text{cachetiming-program } 9c \rangle \equiv$ (12a) 9d \triangleright
 _start:

Uses **_start** 5.

Now we start with initialising the **data** area with some random data. For this we load RDI with the address of the **data** area.

9d $\langle \text{cachetiming-program } 9c \rangle + \equiv$ (12a) $\langle 9c \ 9e \triangleright$
 mov RDI,**data**

Uses **data** 8b.

Next we load the number of bytes to fill into RSI. For this we load the **pagesize** into RSI.

9e $\langle \text{cachetiming-program } 9c \rangle + \equiv$ (12a) $\langle 9d \ 9f \triangleright$
 mov RSI,**pagesize**

Uses **pagesize** 5.

At last we load EDX with some random seed. For this we use **rdtsc** and only use the lower 32 bit of the value.

9f $\langle \text{cachetiming-program } 9c \rangle + \equiv$ (12a) $\langle 9e \ 9g \triangleright$
 rdtsc
 mov EDX,EAX

Now we call **_xorshift** to fill the **data** area.

9g $\langle \text{cachetiming-program } 9c \rangle + \equiv$ (12a) $\langle 9f \ 9h \triangleright$
 call **_xorshift**

Uses **_xorshift** 18a.

2.3.2 Measure Time

Now that we have setup our **data** area we can now cache data from the first page by loading it into a register which also loads this into the cache.

For this we load RDI with the address of the **data** area.

9h $\langle \text{cachetiming-program } 9c \rangle + \equiv$ (12a) $\langle 9g \ 10a \triangleright$
 mov RDI,**data**

Uses **data** 8b.

2 Cache Access Timing

Before we load the data into a register now we will clear the cache lines with the given address. For this we use the instruction `clflush`. After flushing the cache line we ensure (with `lfence`) that all reads from memory are finished before we load the data into a register again (and filling the cache).

```
10a  <cachetiming-program 9c>+≡ (12a) <9h 10b>
      clflush    [RDI]
      lfence
      mov        RCX, [RDI]
```

Now we can determine the time that is needed to load this data once again. We do not need to load RDI again because it has not changed.

```
10b  <cachetiming-program 9c>+≡ (12a) <10a 10d>
      call       _calccachetime
Uses _calccachetime 8a.
```

Now we have the relative cache access time in register RAX. We store this value to the stack and print out an explaining text.

For this we define the text to print and (as a helper) a `carriage return (CR)`.

```
10c  <cachetiming-rodata 10c>≡ (12a) 11b>
      scr:                db 0x0a
      scached:            db "Cached Access Time: ",0x00
```

Defines:

`scached`, used in chunk 10d.
`scr`, used in chunks 10e and 11c.

Now we can print the text.

```
10d  <cachetiming-program 9c>+≡ (12a) <10b 10e>
      push       RAX
      mov        RDI,scached
      call       _print
Uses _print 20a and scached 10c.
```

Then we restore the value and print the measured time to `stdout`. At last we append a `CR` to the output.

```
10e  <cachetiming-program 9c>+≡ (12a) <10d 11a>
      pop        RDI
      mov        RSI,scratch
      call       _printdu64bit
      mov        RSI,scr
      mov        RDI,1
      call       _nprint
```

Uses `_nprint` 19b, `_printdu64bit` 21a, `scr` 10c, and `scratch` 9a.

Now we do the same with an uncached value. The difference is that we do not load the value before.

11a $\langle \text{cachetiming-program } 9c \rangle + \equiv$ (12a) $\langle 10e \ 11c \rangle$

```

    mov     RDI,data
    clflush [RDI]
    lfence
    call    _calccachetime

```

Uses `_calccachetime` 8a and `data` 8b.

Now we have the time of the uncached data access in `RAX` and can print it out with some explaining text.

11b $\langle \text{cachetiming-rodata } 10c \rangle + \equiv$ (12a) $\langle 10c$

```

    suncached:    db "Uncached Access Time: ",0x00

```

Defines:

`suncached`, used in chunk 11c.

11c $\langle \text{cachetiming-program } 9c \rangle + \equiv$ (12a) $\langle 11a \ 11d \rangle$

```

    push     RAX
    mov     RDI,suncached
    call    _print
    pop     RDI
    mov     RSI,scratch
    call    _printdu64bit
    mov     RSI,scr
    mov     RDI,1
    call    _nprint

```

Uses `_nprint` 19b, `_print` 20a, `_printdu64bit` 21a, `scr` 10c, `scratch` 9a, and `suncached` 11b.

At last we exit the program.

11d $\langle \text{cachetiming-program } 9c \rangle + \equiv$ (12a) $\langle 11c$

```

     $\langle \text{exitProgram } 17b \rangle$ 

```

2 Cache Access Timing

Now we can put everything together and have our `cachetiming` program that we can now execute.

```
12a <cachetiming.asm 12a>≡
    <preamble 5>

    section .rodata
    <cachetiming-rodata 10c>

    section .bss
    <cachetiming-udata 8b>

    section .text
    <cachetiming-program 9c>

    <calculate-cache-access-time 8a>

    <xorshift-prng 18a>

    <utilities 17a>
```

The program is placed in `asm/`. With `make` in the folder we can create an executable which is moved to `bin/`. There we can execute this program.

```
$ ./cachetiming
Cached Access Time: 72
Uncached Access Time: 372
$
```

2.4 Read Array via Cache Access Time

Now that we have seen that we can determine if a value was in the cache or not (see [2.3 Measure Cache Access Time](#)) we will read a complete array of data by only measuring the cache access time.

For this we start with some data area that we can read later.

```
12b <cacheread-preamble 12b>≡ (14)
    datasize equ 1024
```

```
12c <cacheread-udata 12c>≡ (14) 13b▷
    align pagesize
    data: resb datasize
```

Uses `data 8b` and `pagesize 5`.

So start with the program and fill the `data` area with some random data (similar to the chunks 9c, 9d, 9e, 9f and 9g).

13a $\langle \text{cacheread-program } 13a \rangle \equiv$ (14) 13c \triangleright

```

_start:
    mov     RDI,data
    mov     RSI,pagesize
    rdtsc
    mov     EDX,EAX
    call    _xorshift

```

Uses `_start` 5, `_xorshift` 18a, `data` 8b, and `pagesize` 5.

Next we will create a probe area that is $256 * \text{pagesize}$. We only access the first byte of each page but we divide the data over such a large area (1 MiB) to ensure that the cache lines that we use do not interfere each other.

13b $\langle \text{cacheread-udata } 12c \rangle + \equiv$ (14) $\triangleleft 12c$

```

probe:          times 256 resb datasize

```

Next we fill this area also with some random data.

13c $\langle \text{cacheread-program } 13a \rangle + \equiv$ (14) $\triangleleft 13a$

```

mov     RDI,probe
mov     RAX,pagesize
mov     RCX,256
mul     RCX
mov     RSI,RCX
rdtsc
mov     EDX,EAX
call    _xorshift

```

Uses `_xorshift` 18a and `pagesize` 5.

2 Cache Access Timing

TBD

14 $\langle \text{cacheread.asm } 14 \rangle \equiv$
 $\langle \text{preamble } 5 \rangle$
 $\langle \text{cacheread-preamble } 12b \rangle$

section .bss
 $\langle \text{cacheread-udata } 12c \rangle$

section .text
 $\langle \text{cacheread-program } 13a \rangle$

$\langle \text{exitProgram } 17b \rangle$

$\langle \text{calculate-cache-access-time } 8a \rangle$

$\langle \text{xorshift-prng } 18a \rangle$

$\langle \text{utilities } 17a \rangle$

3 Signals

3.1 Basics

TBD

3.2 Detecting Signals

TBD

3.3 Handling Signals

TBD

4 Utilities

4.1 Introduction

TBD

17a $\langle utilities\ 17a \rangle \equiv$ (12a 14)

$\langle nprint\ 19b \rangle$

$\langle print\ 20a \rangle$

$\langle printdu64bit\ 21a \rangle$

4.2 Exit Program

TBD

17b $\langle exitProgram\ 17b \rangle \equiv$ (11d 14)

`xor RDI,RDI`

`mov RAX,60`

`syscall`

4.3 Random Number Generator

To initialize the data a [random number generator \(RNG\)](#) is used. The sample programs use [xorshift](#)¹ as [RNG](#).

First we clear the direction flag to ensure that we are incrementing the data pointer RDI.

Next we move the number of values to be generated to RCX (which is a counter in [x86](#) processors) and divide it by 4 (because we use a 32bit [RNG](#)). Additionally we move the seed to EAX.

Parameters

RDI	the address of the memory which is to be filled with random numbers
RSI	the number of bytes that are filled with random numbers. This must be a multiple of 4
EDX	the seed of the RNG

```

18a  <xorshift-prng 18a>≡ (12a 14) 18b>
    _xorshift:
        cld
        mov     RCX,RSI
        shr     RCX,2
        mov     EAX,EDX

```

Defines:

_xorshift, used in chunks 9g and 13.

Now we can generate the next 32bit random number.

```

18b  <xorshift-prng 18a>+≡ (12a 14) <18a 19a>
    .next_random:
        mov     EBX,EAX
        shl     EAX,13
        xor     EAX,EBX
        mov     EBX,EAX
        shr     EAX,17
        xor     EAX,EBX
        mov     EBX,EAX
        shl     EAX,5
        xor     EAX,EBX

```

¹<https://en.wikipedia.org/wiki/Xorshift>

Because we want to generate multiple random numbers we store the value of `EAX` to `[RDI]` and loop for the next random number.

19a `<xorshift-prng 18a>+≡` (12a 14) <18b

```

    stosd
    loop    .next_random
    ret

```

4.4 Printing Strings

4.4.1 Printing Strings with Length

The routine `_nprint` prints a string with the given length to `stdout`.

We move the number of bytes to print to `RDX` which is the 3rd parameter to the systemcall. Next we move the address of the bytes to print to `RSI` which is the 2nd parameter to the systemcall. The 1st argument (in `RDI`) to the systemcall is the file descriptor (1 is `stdout`). Additionally the number of the systemcall (1) is passed in `RAX`. The systemcall (`syscall`) now prints `RDX` bytes from `[RSI]` to the file descriptor `RDI`.

At the end we return to the caller.

Parameters

`RDI` the number of bytes to print to `stdout`

`RSI` the address to the bytes to print to `stdout`

19b `<nprint 19b>≡` (17a)

```

_nprint:
    mov     RDX,RDI
    mov     RDI,1
    mov     RAX,1
    syscall
    ret

```

Defines:

`_nprint`, used in chunks 10e, 11c, 20d, and 23b.

4.4.2 Printing C-Strings

The routine `_print` prints a null-terminated string to `stdout`.

First we clear the direction flag to increment the address in `RDI` while scanning the data.

Next we start with clearing `AL` (setting it to null) and saving the address of the string to `RSI`. We're using `RSI` because we later need the address to calculate the length of the string.

Parameters

RDI the address to the null-terminated bytes to print to **stdout**

20a $\langle print\ 20a \rangle \equiv$ (17a) 20b \triangleright

```

    _print:
        cld
        xor     AL,AL
        mov     RSI,RDI

```

Defines:

 _print, used in chunks 10d and 11c.

Next we search for the terminating **null** (`'\0'`) character. For this we use the instruction **scasb** (scan string byte) which compares the byte at the address **[RDI]** with the value in **AL** and sets the flags accordingly. When the byte at **[RDI]** is not the value of **AL** the next instruction (**jne**) jumps to the given label (**.next_char** in this case).

scasb additionally increments **RDI** so that we go through the string until `'\0'` is found.

20b $\langle print\ 20a \rangle + \equiv$ (17a) \triangleleft 20a 20c \triangleright

```

    .next_char:
        scasb
        jne     .next_char

```

After we have found the string termination we calculate the number of bytes that the string has. In **RSI** we now have the starting address of the bytes to print and in **RDI** we have the end address of the bytes to print. After that we calculate the number of bytes to print.

20c $\langle print\ 20a \rangle + \equiv$ (17a) \triangleleft 20b 20d \triangleright

```

        sub     RDI,RSI

```

Now we have the address of the string in **RDI** and the length of the string in **RSI** which are the 1st and 2nd argument in the call of **_nprint**.

20d $\langle print\ 20a \rangle + \equiv$ (17a) \triangleleft 20c

```

        call    _nprint
        ret

```

Uses **_nprint** 19b.

4.5 Printing Numbers

4.5.1 Printing a 64bit Unsigned Integer

The routine `_printdu64bit` print a given number as unsigned decimal number with 64bit to `stdout`.

To print a decimal number we have to divide the number by 10 and get the remainder for printing (from right to left). For this we move the divisor to a register and the dividend to `RAX`. We have to use `RAX` because this is the only register we can use for division.

Additionally we need the address of the scratch area in `RDI` for storing the result. We also save the address of the scratch area to `R8` for later use.

To increment the address during the processing we clear the direction flag.

Parameters

`RDI` the number number to print to `stdout`

`RSI` the address of a scratch area with a size of at least 20 bytes

```
21a  <printdu64bit 21a>≡ (17a) 21b>
    _printdu64bit:
        mov     RAX,RDI
        mov     RDI,RSI
        mov     R8,RDI
        mov     RCX,10
        cld
```

Defines:

`_printdu64bit`, used in chunks 10e and 11c.

Now we define a label to jump back when we see that there are still more digits to print. Then we test `RAX` for 0 and end the processing of the digits.

```
21b  <printdu64bit 21a>+≡ (17a) <21a 21c>
    .next:
        cmp     RAX,0
        je      .done
```

Next we divide `RAX` by `RCX`. For this we have to clear `RDX` because this is the higher value of the dividend. The result is then placed into `RAX` and the remainder into `RDX`.

```
21c  <printdu64bit 21a>+≡ (17a) <21b 22a>
        xor     RDX,RDX
        div     RCX
```

4 Utilities

We now exchange the result and the remainder because we now need the remainder in RAX (or AL) for further processing. Now we can add the [ASCII](#) character '0' to AL and have the correct [ASCII](#) value in AL. Now we can store the [ASCII](#) character to the scratch area.

```
22a  <printdu64bit 21a>+≡ (17a) <21c 22b>
      xchg      RDX,RAX
      add       AL,'0'
      stosb
```

Now we restore RAX (which we saved to RDX) to go into the next round.

```
22b  <printdu64bit 21a>+≡ (17a) <22a 22c>
      mov       RAX,RDX
      jmp       .next
```

Now that we have all the numbers as [ASCII](#) characters we are nearly done. We now have to reverse the number in memory because the number saved at the lowest address is the digit with the least significance.

We now start with checking if we have written any character. If not then we write the [ASCII](#) character '0' into the memory. We use the instruction `stosb` for this to adjust the address in RDI at the same time.

```
22c  <printdu64bit 21a>+≡ (17a) <22b 22d>
      .done:
      cmp       RDI,RSI
      jne       .printout
      mov       AL,'0'
      stosb
      .printout:
```

Next we calculate the number of digits that the number has. For this we move the address of the last digit to RDX and subtract the start of the scratch area from this. Next we adjust RDI because it points to the first address after the number.

```
22d  <printdu64bit 21a>+≡ (17a) <22c 23a>
      mov       RDX,RDI
      sub       RDX,RSI
      dec       RDI
```

We now have `RSI` with the address of the start of the number and `RDI` with the address of the end. We now have to exchange the digits from the front and the end to get the right number. For this we increment `RSI` and decrement `RDI` after each exchange and when the addresses pass each other we are done.

```
23a  <printdu64bit 21a>+≡ (17a) <22d 23b>
      .reverse:
      mov     AL,[RSI]
      mov     AH,[RDI]
      mov     [RSI],AH
      mov     [RDI],AL
      dec     RDI
      inc     RSI
      cmp     RSI,RDI
      jb      .reverse
```

Now we restore the address of the scratch area to `RSI` and move the number of digits (which we stored in `RDY`) to `RDI` and can the call `_nprint` to print the number.

```
23b  <printdu64bit 21a>+≡ (17a) <23a
      mov     RSI,R8
      mov     RDI,RDX
      call    _nprint
      ret
```

Uses `_nprint` 19b.

A Glossary

x86 x86 denotes a microprocessor architecture based on the 8086/8088 [18](#)

B Acronyms

ASCII American Standard Code for Information Interchange [22](#)

CR carriage return [10](#)

RNG random number generator [18](#)

C x86-Instructions

`clflush` Flush Cache Line, introduced with Intel[®] Pentium[®] 4 [10](#)

`lfence` Load Fence, introduced with Intel[®] Pentium[®] 4 [8](#), [10](#)

`rdtsc` Read Time Stamp Counter, introduced with Intel[®] Pentium[®] [7](#), [9](#)

D Code Chunks

<cacheread-preamble 12b>
<cacheread-program 13a>
<cacheread-udata 12c>
<cacheread.asm 14>
<cachetiming-program 9c>
<cachetiming-rodata 10c>
<cachetiming-udata 8b>
<cachetiming.asm 12a>
<calculate-cache-access-time 8a>
<exitProgram 17b>
<nprint 19b>
<preamble 5>
<print 20a>
<printdu64bit 21a>
<scratch-data 9a>
<tsc-64bit 7>
<utilities 17a>
<xorshift-prng 18a>