

Meltdown and Spectre Samples

Written in Assembly

U. Plonus

March 10, 2018

Contents

1	Introduction	5
1.1	Overview	5
1.2	Nasm	5
2	Cache Access Timing	7
2.1	Introduction	7
2.2	Detect Cache Access Time	7
2.2.1	High Resolution Timer	7
2.2.2	Cache Access Time Routine	8
2.3	Measure Cache Access Time	8
2.4	Read Array via Cache Access Time	10
3	Signals	11
3.1	Basics	11
3.2	Detecting Signals	11
3.3	Handling Signals	11
4	Utilities	13
4.1	Introduction	13
4.2	Exit Program	13
4.3	Random Number Generator	14
4.4	Printing Strings	15
4.4.1	Printing Strings with Length	15
4.4.2	Printing C-Strings	15
4.5	Printing Numbers	16
A	Glossary	17
B	Acronyms	19
C	x86-Instructions	21
D	Code Chunks	23

1 Introduction

1.1 Overview

TBD

1.2 Nasm

TBD

```
5  <preamble 5>≡ (9f)
    bits 64
        global _start
        pagesize equ 4096
```

Defines:

`_start`, used in chunk 9a.

`pagesize`, used in chunks 8b and 9c.

2 Cache Access Timing

2.1 Introduction

TBD

2.2 Detect Cache Access Time

2.2.1 High Resolution Timer

First we need a high resolution timer to determine the cache access time. For this Intel[®] processors since the Pentium[®] model have a time stamp counter. The time stamp counter is monotonically incrementing. When reading the time stamp counter (with `rdtsc`) the result is delivered back in the registers `EDX` and `EAX` forming a 64bit value. The time stamp counter is not an absolute value but a relative value, meaning that you cannot (easily) calculate from the time stamp counter to some time units (e.g. ns). But this is no problem as we only want to measure relative times.

To retrieve a 64bit value for the time we shift the value in `EDX` 32 bits to the left and add the value of `EAX` to this.

7 $\langle tsc-64bit \ 7 \rangle \equiv$ (8a)
 `rdtsc`
 `shl RDX,32`
 `add RAX,RDX`

2.2.2 Cache Access Time Routine

Next we need a routine that calculates the cache access time for us.

First we have to ensure in this routine that the speculative execution of the processor does not interfere with our time measurement. For this we use the instruction `lfence` which ensures that all previous reads are done before executing the next instructions.

Next we access a memory location with the address `RDI` by loading this into `RCX` and measure the time before and after the access.

The command `lfence` before reading the time stamp counter is needed because we have to ensure that all reads before the time measurements are done.

At last we calculate the relative time needed to access the memory location. In theory we should see a difference whether the memory location is accessed before or not.

Parameters

`RDI` the address of the memory which is loaded either from the cache or from memory

```
8a  <calculate-cache-access-time 8a>≡ (10a)
    _calccachetime:
        lfence
        <tsc-64bit 7>
        mov     R8,RAX
        mov     RCX,[RDI]
        lfence
        <tsc-64bit 7>
        sub     RAX,R8
        ret
```

Defines:

`_calccachetime`, never used.

2.3 Measure Cache Access Time

To measure the cache timing we create a standalone program that shows us the time for a cached and for an uncached memory access.

First we need some area in memory with data which we can later read from. This data area goes into the area `.bss` which contains uninitialized data. We align the data at a page boundary and reserve 2 pages for our data.

```
8b  <cachetiming-uninitialized-data 8b>≡ (9f)
    section .bss
        align    pagesize
        data:    times 2 resb pagesize
```

Defines:

`data`, used in chunk 9b.

Uses `pagesize` 5.

The program begins with the label `_start` and is in the section `.text`.

```
9a  <cachetiming-program 9a>≡ (9f) 9b>
    section .text
    _start:
Uses _start 5.
```

Now we start with initialising the `data` area with some random data. For this we load `RDI` with the address of the `data` area.

```
9b  <cachetiming-program 9a>+≡ (9f) <9a 9c>
    mov     RDI,data
Uses data 8b.
```

Next we load the number of bytes to fill into `RSI`. For this we load the `pagesize` into `RSI` and multiply it with 2 by shifting the value 1 bit to the left.

```
9c  <cachetiming-program 9a>+≡ (9f) <9b 9d>
    mov     RSI,pagesize
    shl     RSI,1
Uses pagesize 5.
```

At last we load `EDX` with some random seed. For this we use `rdtsc` and only use the lower 32 bit of the value.

```
9d  <cachetiming-program 9a>+≡ (9f) <9c 9e>
    rdtsc
    mov     EDX,EAX
```

Now we call `_xorshift` to fill the `data` area.

```
9e  <cachetiming-program 9a>+≡ (9f) <9d 10a>
    call    _xorshift
Uses _xorshift 14a.
```

TBD

```
9f  <cachetiming.asm 9f>≡
    <preamble 5>

    <cachetiming-rodata 10b>

    <cachetiming-uninitialized-data 8b>

    <cachetiming-program 9a>
```

2 Cache Access Timing

10a $\langle \text{cachetiming-program } 9a \rangle + \equiv$ (9f) <9e

```
    mov     RDI,scached
    call    _print
    mov     RDI,scr
    mov     RSI,1
    call    _nprint
    mov     RDI,suncached
    call    _print
    mov     RDI,scr
    mov     RSI,1
    call    _nprint
     $\langle \text{exitProgram } 13b \rangle$ 

     $\langle \text{calculate-cache-access-time } 8a \rangle$ 

     $\langle \text{xorshift-prng } 14a \rangle$ 

     $\langle \text{utilities } 13a \rangle$ 
    Uses _nprint 15a, _print 15b, scached 10b, scr 10b, and suncached 10b.
    TBD

10b  $\langle \text{cachetiming-rodata } 10b \rangle \equiv$  (9f)



```
 section .rodata
 suncached: db "Uncached Access Time: ",0x00
 scached: db "Cached Access Time: ",0x00
 scr: db 0x0a
```



Defines:



- scached, used in chunk 10a.
- scr, used in chunk 10a.
- suncached, used in chunk 10a.

```

2.4 Read Array via Cache Access Time

TBD

3 Signals

3.1 Basics

TBD

3.2 Detecting Signals

TBD

3.3 Handling Signals

TBD

4 Utilities

4.1 Introduction

TBD

$$\begin{aligned} 13a \quad \langle utilities \ 13a \rangle &\equiv & (10a) \\ &\langle nprint \ 15a \rangle \\ &\langle print \ 15b \rangle \end{aligned}$$

4.2 Exit Program

TBD

$$\begin{aligned} 13b \quad \langle exitProgram \ 13b \rangle &\equiv & (10a) \\ &\quad xor \quad RDI, RDI \\ &\quad mov \quad RAX, 60 \\ &\quad syscall \end{aligned}$$

4.3 Random Number Generator

To initialize the data a [random number generator \(RNG\)](#) is used. The sample programs use [xorshift](#)¹ as [RNG](#).

First we move the number of values to be generated to **RCX** (which is a counter in [x86](#) processors) and divide it by 4 (because we use a 32bit [RNG](#)). Additionally we move the seed to **EAX**.

Parameters

RDI	the address of the memory which is to be filled with random numbers
RSI	the number of bytes that are filled with random numbers. This must be a multiple of 4
EDX	the seed of the RNG

14a $\langle \text{xorshift-prng } 14a \rangle \equiv$ (10a) 14b \triangleright

```

    _xorshift:
        mov     RCX,RSI
        shr     RCX,2
        mov     EAX,EDX

```

Defines:

 _xorshift, used in chunk 9e.

Now we can generate the next 32bit random number.

14b $\langle \text{xorshift-prng } 14a \rangle + \equiv$ (10a) \triangleleft 14a 14c \triangleright

```

    .next_random:
        mov     EBX,EAX
        shl     EAX,13
        xor     EAX,EBX
        mov     EBX,EAX
        shr     EAX,17
        xor     EAX,EBX
        mov     EBX,EAX
        shl     EAX,5
        xor     EAX,EBX

```

Because we want to generate multiple random numbers we store the value of **EAX** to **[RDI]** and loop for the next random number.

14c $\langle \text{xorshift-prng } 14a \rangle + \equiv$ (10a) \triangleleft 14b

```

        stosd
        loop    .next_random
        ret

```

¹<https://en.wikipedia.org/wiki/Xorshift>

4.4 Printing Strings

4.4.1 Printing Strings with Length

The routine `_nprint` prints a string with the given length to `stdout`.

We move the number of bytes to print to `RDX` which is the 3rd parameter to the `syscall`. Next we move the address of the bytes to print to `RSI` which is the 2nd parameter to the `syscall`. The 1st argument (in `RDI`) to the `syscall` is the file descriptor (1 is `stdout`). Additionally the number of the `syscall` (1) is passed in `RAX`. The `syscall` (`syscall`) now prints `RDX` bytes from `[RSI]` to the file descriptor `RDI`.

At the end we return to the caller.

Parameters

`RDI` the number of bytes to print to `stdout`

`RSI` the address to the bytes to print to `stdout`

15a $\langle nprint\ 15a \rangle \equiv$ (13a)
 `_nprint:`
 `mov` `RDY,RSI`
 `mov` `RSI,RDI`
 `mov` `RDI,1`
 `mov` `RAX,1`
 `syscall`
 `ret`

Defines:

`_nprint`, used in chunks 10a and 16c.

4.4.2 Printing C-Strings

The routine `_print` prints a null-terminated string to `stdout`.

So first we start with clearing `AL` (setting it to null) and saving the address of the string to `RSI`. We're using `RSI` because we later need the address to calculate the length of the string.

Parameters

`RDI` the address to the null-terminated bytes to print to `stdout`

15b $\langle print\ 15b \rangle \equiv$ (13a) 16a▷
 `_print:`
 `xor` `AL,AL`
 `mov` `RSI,RDI`

Defines:

`_print`, used in chunk 10a.

4 Utilities

Next we search for the terminating `null` (`'\0'`) character. For this we use the instruction `scasb` (scan string byte) which compares the byte at the address `[RDI]` with the value in `AL` and sets the flags accordingly. When the byte at `[RDI]` is not the value of `AL` the next instruction (`jne`) jumps to the given label (`.next_char` in this case).

`scasb` additionally increments `RDI` so that we go through the string until `'\0'` is found.

```
16a  <print 15b>+≡ (13a) <15b 16b>
      .next_char:
      scasb
      jne     .next_char
```

After we have found the string termination we calculate the number of bytes that the string has. For this we exchange the registers `RDI` and `RSI`. In `RDI` we now have the starting address of the bytes to print and in `RSI` we have the end address of the bytes to print. After that we calculate the number of bytes to print.

```
16b  <print 15b>+≡ (13a) <16a 16c>
      xchg     RDI,RSI
      sub      RSI,RDI
```

Now we have the address of the string in `RDI` and the length of the string in `RSI` which are the 1st and 2nd argument in the call of `_nprint`.

```
16c  <print 15b>+≡ (13a) <16b
      call     _nprint
      ret
```

Uses `_nprint 15a`.

4.5 Printing Numbers

TBD

A Glossary

x86 x86 denotes a microprocessor architecture based on the 8086/8088 [14](#)

B Acronyms

RNG random number generator [14](#)

C x86-Instructions

`lfence` Load Fence [8](#)

`rdtsc` Read Time Stamp Counter [7](#), [9](#)

D Code Chunks

⟨cachetiming-program 9a⟩
⟨cachetiming-rodata 10b⟩
⟨cachetiming-uninitialized-data 8b⟩
⟨cachetiming.asm 9f⟩
⟨calculate-cache-access-time 8a⟩
⟨exitProgram 13b⟩
⟨nprint 15a⟩
⟨preamble 5⟩
⟨print 15b⟩
⟨tsc-64bit 7⟩
⟨utilities 13a⟩
⟨xorshift-prng 14a⟩