

Meltdown and Spectre Samples

Written in Assembly

U. Plonus

March 10, 2018

Contents

1	Introduction	5
1.1	Overview	5
1.2	Nasm	5
2	Cache Access Timing	7
2.1	Introduction	7
2.2	Detect Cache Access Time	7
2.2.1	High Resolution Timer	7
2.2.2	Cache Access Time Routine	8
2.3	Measure Cache Access Time	9
2.3.1	Setup	9
2.3.2	Measure Time	10
2.4	Read Array via Cache Access Time	11
3	Signals	13
3.1	Basics	13
3.2	Detecting Signals	13
3.3	Handling Signals	13
4	Utilities	15
4.1	Introduction	15
4.2	Exit Program	15
4.3	Random Number Generator	16
4.4	Printing Strings	17
4.4.1	Printing Strings with Length	17
4.4.2	Printing C-Strings	17
4.5	Printing Numbers	18
A	Glossary	19
B	Acronyms	21
C	x86-Instructions	23
D	Code Chunks	25

1 Introduction

1.1 Overview

TBD

1.2 Nasm

TBD

```
5  <preamble 5>≡ (10d)
    bits 64
        global _start
        pagesize equ 4096
```

Defines:

 _start, used in chunk 9b.
 pagesize, used in chunk 9.

2 Cache Access Timing

2.1 Introduction

TBD

2.2 Detect Cache Access Time

2.2.1 High Resolution Timer

First we need a high resolution timer to determine the cache access time. For this we use the time stamp counter. The time stamp counter is monotonically incrementing. When reading the time stamp counter (with `rdtsc`) the result is delivered back in the registers `EDX` and `EAX` forming a 64bit value. The time stamp counter is not an absolute value but a relative value, meaning that you cannot (easily) calculate from the time stamp counter to some time units (e.g. ns). But this is no problem as we only want to measure relative times.

To retrieve a 64bit value for the time we shift the value in `EDX` 32 bits to the left and add the value of `EAX` to this.

7 $\langle tsc-64bit \ 7 \rangle \equiv$ (8)

```
rdtsc
shl    RDX,32
add    RAX,RDX
```

2.2.2 Cache Access Time Routine

Next we need a routine that calculates the cache access time for us.

First we have to ensure in this routine that the speculative execution of the processor does not interfere with our time measurement. For this we use the instruction `lfence` which ensures that all previous reads are done before executing the next instructions.

Next we access a memory location with the address `RDI` by loading this into `RCX` and measure the time before and after the access.

The command `lfence` before reading the time stamp counter is needed because we have to ensure that all reads before the time measurements are done.

At last we calculate the relative time needed to access the memory location. In theory we should see a difference whether the memory location is accessed before or not.

Parameters

`RDI` the address of the memory which is loaded either from the cache or from memory

```

8  <calculate-cache-access-time 8>≡ (11a)
    _calccachetime:
        lfence
        <tsc-64bit 7>
        mov     R8,RAX
        mov     RCX,[RDI]
        lfence
        <tsc-64bit 7>
        sub     RAX,R8
        ret

```

Defines:

`_calccachetime`, used in chunk 10c.

2.3 Measure Cache Access Time

2.3.1 Setup

To measure the cache timing we create a standalone program that shows us the time for a cached and for an uncached memory access.

First we need some area in memory with data which we can later read from. This data area goes into the area `.bss` which contains uninitialized data. We align the data at a page boundary and reserve 2 pages for our data.

```
9a  <cachetiming-uninitialized-data 9a>≡ (10d)
    section .bss
        align      pagesize
        data:      times 2 resb pagesize
```

Defines:

`data`, used in chunks 9c and 10a.

Uses `pagesize` 5.

The program begins with the label `_start` and is in the section `.text`.

```
9b  <cachetiming-program 9b>≡ (10d) 9c>
    section .text
    _start:
```

Uses `_start` 5.

Now we start with initialising the `data` area with some random data. For this we load RDI with the address of the `data` area.

```
9c  <cachetiming-program 9b>+≡ (10d) <9b 9d>
        mov     RDI,data
```

Uses `data` 9a.

Next we load the number of bytes to fill into RSI. For this we load the `pagesize` into RSI and multiply it with 2 by shifting the value 1 bit to the left.

```
9d  <cachetiming-program 9b>+≡ (10d) <9c 9e>
        mov     RSI,pagesize
        shl     RSI,1
```

Uses `pagesize` 5.

At last we load EDX with some random seed. For this we use `rdtsc` and only use the lower 32 bit of the value.

```
9e  <cachetiming-program 9b>+≡ (10d) <9d 9f>
        rdtsc
        mov     EDX,EAX
```

Now we call `_xorshift` to fill the `data` area.

```
9f  <cachetiming-program 9b>+≡ (10d) <9e 10a>
        call    _xorshift
```

Uses `_xorshift` 16a.

2 Cache Access Timing

2.3.2 Measure Time

Now that we have setup our `data` area we can now cache data from the first page by loading it into a register which also loads this into the cache.

For this we load `RDI` with the address of the `data` area.

```
10a  <cachetiming-program 9b>+≡ (10d) <9f 10b>
      mov      RDI,data
```

Uses `data 9a`.

Before we load the data into a register now we will clear the cache lines with the given address. For this we use the instruction `clflush`. After that we will load the data into a register.

```
10b  <cachetiming-program 9b>+≡ (10d) <10a 10c>
      clflush   [RDI]
      mov      RCX,[RDI]
```

Now we can determine the time that is needed to load this data once again. We do not need to load `RDI` again because it has not changed.

```
10c  <cachetiming-program 9b>+≡ (10d) <10b 11a>
      call     _calccachetime
```

Uses `_calccachetime 8`.

Now we have the relative cache access time in register `RAX`. Now we can print the measured time to `stdout`.

TBD

```
10d  <cachetiming.asm 10d>≡
      <preamble 5>

      <cachetiming-rodata 11b>

      <cachetiming-uninitialized-data 9a>

      <cachetiming-program 9b>
```

11a $\langle \text{cachetiming-program } 9b \rangle + \equiv$ (10d) $\triangleleft 10c$

```

    mov     RDI,scached
    call    _print
    mov     RDI,scr
    mov     RSI,1
    call    _nprint
    mov     RDI,suncached
    call    _print
    mov     RDI,scr
    mov     RSI,1
    call    _nprint
     $\langle \text{exitProgram } 15b \rangle$ 

```

$\langle \text{calculate-cache-access-time } 8 \rangle$

$\langle \text{xorshift-prng } 16a \rangle$

$\langle \text{utilities } 15a \rangle$

Uses `_nprint` 17b, `_print` 18a, `scached` 11b, `scr` 11b, and `suncached` 11b.

TBD

11b $\langle \text{cachetiming-rodata } 11b \rangle \equiv$ (10d)

```

    section .rodata
        suncached:    db "Uncached Access Time: ",0x00
        scached:      db "Cached Access Time: ",0x00
        scr:           db 0x0a

```

Defines:

`scached`, used in chunk 11a.

`scr`, used in chunk 11a.

`suncached`, used in chunk 11a.

2.4 Read Array via Cache Access Time

TBD

3 Signals

3.1 Basics

TBD

3.2 Detecting Signals

TBD

3.3 Handling Signals

TBD

4 Utilities

4.1 Introduction

TBD

$$\begin{aligned} 15a \quad \langle utilities \ 15a \rangle &\equiv & (11a) \\ &\langle nprint \ 17b \rangle \\ &\langle print \ 18a \rangle \end{aligned}$$

4.2 Exit Program

TBD

$$\begin{aligned} 15b \quad \langle exitProgram \ 15b \rangle &\equiv & (11a) \\ &\quad xor \quad RDI, RDI \\ &\quad mov \quad RAX, 60 \\ &\quad syscall \end{aligned}$$

4.3 Random Number Generator

To initialize the data a [random number generator \(RNG\)](#) is used. The sample programs use [xorshift](#)¹ as [RNG](#).

First we clear the direction flag to ensure that we are incrementing the data pointer RDI.

Next we move the number of values to be generated to RCX (which is a counter in [x86](#) processors) and divide it by 4 (because we use a 32bit [RNG](#)). Additionally we move the seed to EAX.

Parameters

RDI	the address of the memory which is to be filled with random numbers
RSI	the number of bytes that are filled with random numbers. This must be a multiple of 4
EDX	the seed of the RNG

```

16a  <xorshift-prng 16a>≡ (11a) 16b>
      _xorshift:
          cld
          mov     RCX,RSI
          shr     RCX,2
          mov     EAX,EDX

```

Defines:

_xorshift, used in chunk 9f.

Now we can generate the next 32bit random number.

```

16b  <xorshift-prng 16a>+≡ (11a) <16a 17a>
      .next_random:
          mov     EBX,EAX
          shl     EAX,13
          xor     EAX,EBX
          mov     EBX,EAX
          shr     EAX,17
          xor     EAX,EBX
          mov     EBX,EAX
          shl     EAX,5
          xor     EAX,EBX

```

¹<https://en.wikipedia.org/wiki/Xorshift>

Because we want to generate multiple random numbers we store the value of `EAX` to `[RDI]` and loop for the next random number.

```
17a  <xorshift-prng 16a>+≡ (11a) <16b
      stosd
      loop    .next_random
      ret
```

4.4 Printing Strings

4.4.1 Printing Strings with Length

The routine `_nprint` prints a string with the given length to `stdout`.

We move the number of bytes to print to `RDX` which is the 3rd parameter to the systemcall. Next we move the address of the bytes to print to `RSI` which is the 2nd parameter to the systemcall. The 1st argument (in `RDI`) to the systemcall is the file descriptor (1 is `stdout`). Additionally the number of the systemcall (1) is passed in `RAX`. The systemcall (`syscall`) now prints `RDX` bytes from `[RSI]` to the file descriptor `RDI`.

At the end we return to the caller.

Parameters

`RDI` the number of bytes to print to `stdout`

`RSI` the address to the bytes to print to `stdout`

```
17b  <nprint 17b>≡ (15a)
      _nprint:
      mov     RDX,RSI
      mov     RSI,RDI
      mov     RDI,1
      mov     RAX,1
      syscall
      ret
```

Defines:

`_nprint`, used in chunks 11a and 18d.

4.4.2 Printing C-Strings

The routine `_print` prints a null-terminated string to `stdout`.

First we clear the direction flag to increment the address in `RDI` while scanning the data.

Next we start with clearing `AL` (setting it to null) and saving the address of the string to `RSI`. We're using `RSI` because we later need the address to calculate the length of the string.

Parameters

RDI the address to the null-terminated bytes to print to `stdout`

```
18a  <print 18a>≡ (15a) 18b>
      _print:
          cld
          xor     AL,AL
          mov     RSI,RDI
```

Defines:

`_print`, used in chunk 11a.

Next we search for the terminating `null` (`'\0'`) character. For this we use the instruction `scasb` (scan string byte) which compares the byte at the address `[RDI]` with the value in `AL` and sets the flags accordingly. When the byte at `[RDI]` is not the value of `AL` the next instruction (`jne`) jumps to the given label (`.next_char` in this case).

`scasb` additionally increments `RDI` so that we go through the string until `'\0'` is found.

```
18b  <print 18a>+≡ (15a) <18a 18c>
      .next_char:
          scasb
          jne     .next_char
```

After we have found the string termination we calculate the number of bytes that the string has. For this we exchange the registers `RDI` and `RSI`. In `RDI` we now have the starting address of the bytes to print and in `RSI` we have the end address of the bytes to print. After that we calculate the number of bytes to print.

```
18c  <print 18a>+≡ (15a) <18b 18d>
          xchg     RDI,RSI
          sub      RSI,RDI
```

Now we have the address of the string in `RDI` and the length of the string in `RSI` which are the 1st and 2nd argument in the call of `_nprint`.

```
18d  <print 18a>+≡ (15a) <18c
          call     _nprint
          ret
```

Uses `_nprint` 17b.

4.5 Printing Numbers

TBD

A Glossary

x86 x86 denotes a microprocessor architecture based on the 8086/8088 [16](#)

B Acronyms

RNG random number generator [16](#)

C x86-Instructions

`clflush` Flush Cache Line, introduced with Intel[®] Pentium[®] 4 [10](#)

`lfence` Load Fence, introduced with Intel[®] Pentium[®] 4 [8](#)

`rdtsc` Read Time Stamp Counter, introduced with Intel[®] Pentium[®] [7](#), [9](#)

D Code Chunks

⟨cachetiming-program 9b⟩
⟨cachetiming-rodata 11b⟩
⟨cachetiming-uninitialized-data 9a⟩
⟨cachetiming.asm 10d⟩
⟨calculate-cache-access-time 8⟩
⟨exitProgram 15b⟩
⟨nprint 17b⟩
⟨preamble 5⟩
⟨print 18a⟩
⟨tsc-64bit 7⟩
⟨utilities 15a⟩
⟨xorshift-prng 16a⟩